


SpaceX Falcon 9

First stage Landing Prediction

Applied Machine Learning Pipeline to Forecast Booster Reusability & Launch Economics

Capstone Project Report | Data Science & Predictive Analytics



Outline

1. Executive Summary

2. Introduction

3. Methodology

4. Results

5. Conclusion

Executive Summary

Core Objective

Develop an end-to-end classification pipeline to predict whether the SpaceX Falcon 9 first stage booster will successfully land, enabling competitive cost benchmarking against traditional launch providers.

Key Outcomes

Achieved 83.3% test classification accuracy across evaluated models. Identified key success drivers including payload mass thresholds, orbit dynamics, and launch site geographic parameters.

Introduction

\$62M

SpaceX Commercial Launch Price

\$165M


Traditional launch providers

The Reusability Advantage

Traditional launch providers incur costs upward of **\$165 million** per launch because rockets are single-use.

SpaceX delivers over **60% cost savings** by recovering and refurbishing the Falcon 9 first stage booster, revolutionizing commercial space transportation market dynamics.

Data Science Methodology



Data Collection Methodology



SpaceX REST API

Extracted raw launch details, booster versions, payload specifications, launchpads, and landing outcomes across historical missions.

Pulling SpaceX Launch Data from Wikipedia

- **Grab Falcon 9 and Falcon Heavy launch history** straight from Wikipedia.
- **Pull up the web page** using the link for Falcon 9's launch records.
- **Snag the column headers** from the table to see what data fields we're working with.
- **Turn the table data into a clean DataFrame** so it's ready to analyze.

Data Wrangling Methodology



Web Scraping

Scraped supplementary Falcon 9 historical records from Wikipedia using BeautifulSoup to validate landing pad coordinates and mission payloads.

Cleaning & Preparing the Launch Data

- **Count up launches by location** to see which sites are used the most.
- **Break down orbit types** to see how frequently each target orbit is targeted.
- **Track success rates per orbit** to check how mission outcomes vary depending on the destination.
- **Build a landing success label** using raw data from the main Outcome column.

Predictive Analysis Methodology

- **Extract the target variable:** Grab the 'Class' column as a Pandas Series, convert it to a NumPy array, and store it in Y.
- **Scale the feature set:** Standardize the feature matrix X so all variables are on the same scale, then overwrite X.
- **Split the data:** Divide features and labels into training and testing sets (holding back 20% for testing with `random_state=2` for consistency).
- **Tune Logistic Regression:** Set up Logistic Regression with 10-fold cross-validation (`GridSearchCV`) to test hyperparameters and find the best-performing configuration.

Data Wrangling with SQL

Exploring the Data with SQL Queries

- Find all unique launch pads used across the missions.
- Pull 5 sample records for launch sites starting with "CCA".
- Calculate total payload weight delivered for NASA Commercial Resupply Services (CRS) missions.


Launch_Site
CCAFS LC-40
VAFB SLC-4E
KSC LC-39A
CCAFS SLC-40

Landing_Outcome	Numbers
Success (drone ship)	5
Success (ground pad)	3

- Calculate average payload weight specifically for the F9 v1.1 booster.
- Find the date of the very first successful ground-pad landing.
- Identify boosters that landed on a drone ship while carrying a payload between 4,000 kg and 6,000 kg.
- Count up the overall wins and losses by totaling successful vs. failed mission outcomes.

- Find the heaviest payloads carried and list the booster versions responsible using a subquery.
- Filter for 2015 drone ship failures, pulling up the month, booster version, and launch site for each instance.
- Rank landing successes by frequency between June 4, 2010, and March 20, 2017 (highest to lowest).


EDA with Visualization




Visualised launch success rate throughout the years

EDA with Visualization

Visualised Success rates of
each orbit type



Launch Outcome Proportions by all Sites with Dash



Correlation between payload mass and Launch Success with Dash


4



Booster Version Category

- v1.0
- v1.1
- FT
- B4
- B5

Interactive Visual Analytics with Folium



Folium Map Analytics: Mapped launch site proximities to coastlines, railways, and highways to evaluate safety margins and logistical constraints.

Predictive Models Evaluation

Classification Algorithm	Best Hyperparameters	Train Accuracy	Test Accuracy
Decision Tree Classifier	criterion: entropy, max_depth: 4, splitter: random	84.7%	83.3%
Logistic Regression	C: 0.01, penalty: l2, solver: lbfgs	84.7%	83.3%
Support Vector Machine (SVM)	C: 1.0, kernel: rbf, gamma: scale	84.7%	83.3%
K-Nearest Neighbors (KNN)	n_neighbors: 5, algorithm: auto, p: 2	84.7%	83.3%

All optimized classifiers achieved an identical top test accuracy of 83.33%, demonstrating reliable generalization on unseen test sets.

Key Machine Learning Insights

- ✓ Payload Mass Impact: Light payloads (< 5,300 kg) exhibit significantly higher landing success rates across all launch sites.
- ✓ Flight Count Progression: Booster reliability improved drastically after Flight #20 as block revisions (Block 5) matured.
- ✓ Coastal Proximity: All successful drone ship and land pad recovery sites are strategically placed along coastlines for trajectory safety.
- ✓ Model Generalization: Hyperparameter tuning via GridSearchCV successfully prevented overfitting across all classification algorithms.

Strategic Value & Market Impact

Competitive Bidding

By predicting booster landing probability for target missions, rival launch startups can accurately model SpaceX launch margins and formulate competitive commercial bids.

Reusability Forecasting

Data-driven predictive pipelines offer essential risk management insights for satellite operators deciding between re-flown or new booster stages.

Questions & Discussion

SpaceX Falcon 9 First Stage Landing Prediction Project

IBM Data Science Capstone Presentation | Machine Learning & Predictive Modeling

Conclusion

- Analysis Complete: We've turned raw web pages into a clean database that's ready to power our decision-making.
- Key Patterns: We now have clear visibility into how payload weight, launch sites, and booster types impact landing rates.
- Strategy: We can confidently use these insights and models to price launches better, manage risk, and plan future missions.

Appendix



- <https://github.com/dimi-kaffes/Data-analysis-capstone>



- <https://www.imaging-resource.com/news/how-the-most-prolific-modern-launch-photographer-catches-rockets/>